

EXAGEOSTAT-R INSTALLATION GUIDE

VERSION 1.0.0

November 11, 2018

CONTENTS

1	Chapter 1. General Information	2
1.1	Introduction	2
1.2	Supported Platforms	2
1.3	Software Dependencies	2
2	Chapter 2. macOS	3
2.1	Supported Releases	3
2.2	Installing the prerequisite Packages	3
2.3	ExaGeoStat on R	4
2.4	Verifying the Installation	5
3	Chapter 3. Linux	5
3.1	Supported Linux Distributions	5
3.2	Installing the prerequisite Packages	5
3.3	ExaGeoStat on R	7
3.4	Verifying the Installation	7
4	Chapter 4. Examples	8
4.1	Example 1: Synthetic generation of Geo-spatial data with MLE exact computation,	8
4.2	Example 2: Synthetic generation of Geo-spatial data with MLE TLR computation,	8
4.3	Example 3: Synthetic generation of Geo-spatial data with MLE DST computation,	9
4.4	Example 4: Synthetic generation of measurements based on given Geo-spatial data locations with MLE Exact computation,	9

* Extreme Computing Research Center (ECRC), KAUST, Thuwal, Saudi Arabia

1 CHAPTER 1. GENERAL INFORMATION

1.1 Introduction

This document describes how to install ExaGeoStatR on various platforms. One chapter is dedicated to each operating system.

1.2 Supported Platforms

ExaGeoStatR has been tested and is supported on the following operating systems:

- macOS
- Linux distributions

Note: ExaGeoStat can be run practically on any Linux-like environment with a decent and fairly up-to-date C++ compiler, for example gcc 5.x. Certain ExaGeoStat-R features (Dense, Diagonal Super Tile(DST), Tile Low Rank(TLR), etc.) depend on the availability of external libraries (Nlopt, StarPU, Chameleon, HiCMA, Stars-H, etc.)

1.3 Software Dependencies

1. **Portable Hardware Locality (hwloc):** a software package providing a portable abstraction of the hierarchical topology of modern architectures.
2. **Nlopt: a library for nonlinear optimization** providing a common interface for a number of different optimization routines freely available online as well as original implementations of various other algorithms.
3. **GNU Scientific Library (GSL):** a collection of routines for numerical computing.
4. **StarPU:** a dynamic runtime system library for task-based programming model running on shared/distributed-memory architectures as well as GPU-based systems.
5. **Chameleon:** a dense linear algebra software relying on sequential task-based algorithms and a dynamic runtime system.
6. **Hierarchical Computations on Manycore Architectures library (HiCMA):** a low-rank matrix computation library exploiting the data sparsity of the matrix operator.
7. **Software for Testing Accuracy, Reliability and Scalability of Hierarchical computations (STARS-H):** a high performance low-rank matrix approximation library generating low-rank matrices on shared/distributed-memory systems.

2 CHAPTER 2. MACOS

2.1 Supported Releases

This Chapter provides additional information for installing ExaGeoStat on macOS. The following release is supported:

1. macOS high sierra 10.13.

2.2 Installing the prerequisite Packages

1. Downloading, Configuring and Building CMake 3.2.3 or higher,

```
$ wget https://cmake.org/files/v3.10/cmake-3.10.0-rc3.tar.gz
$ tar -zxvf cmake-3.10.0-rc3.tar.gz
$ cd cmake-3.10.0-rc3
$ ./configure
$ make -j && make -j install
```

2. Download an Install Intel MKL library from <https://software.intel.com/en-us/mkl> for macOS.

```
$ export MKLROOT=/opt/intel/mkl
```

3. Downloading, Configuring and Building NLOPT library 2.4.2 or higher,

```
$ wget http://ab-initio.mit.edu/nlopt/nlopt-2.4.2.tar.gz
$ tar -zxvf nlopt-2.4.2.tar.gz
$ cd nlopt-2.4.2
$ ./configure --enable-shared --without-guile
$ make -j && make -j install
```

4. Downloading, Configuring and Building GSL library 2.4 or higher,

```
$ wget https://ftp.gnu.org/gnu/gsl/gsl-2.4.tar.gz
$ tar -zxvf gsl-2.4.tar.gz
$ cd gsl-2.4
$ ./configure
$ make -j && make -j install
```

5. Downloading, Configuring and Building hwloc 1.11.5 or higher,

```
$ wget https://www.open-mpi.org/software/hwloc/v1.11/downloads/hwloc-1.11.5.tar.gz
$ tar -zxvf hwloc-1.11.5.tar.gz
$ cd hwloc-1.11.5
$ ./configure
$ make -j && make -j install
```

6. Downloading, Configuring and Building StarPU 1.2.5 or higher,

```
$ wget http://starpup.gforge.inria.fr/files/starpup-1.2.5/starpup-1.2.5.tar.gz
$ tar -zxvf starpu-1.2.5.tar.gz
$ cd starpu-1.2.5
$ ./configure --disable-cuda --disable-mpi --disable-opencl
$ make -j && make -j install
```

7. Download Chameleon, HiCMA, and STARS-H,

```
$ git clone https://github.com/ecrc/exageostat.git
$ cd exageostat
$ git submodule update --init --recursive
```

8. Configuring and Building Chameleon Software,

```
$ cd exageostat
$ cd hicma
$ cd chameleon
$ mkdir build && cd build
$ cmake .. -DCHAMELEON_USE_MPI=OFF -DBUILD_SHARED_LIBS=ON
-DCLAS_DIR="${MKLR00T}" -DLAPACK_DIR="${MKLR00T}"
-DBLAS_LIBRARIES="-L${MKLR00T}/lib;-lmkl_intel_lp64;-lmkl_core;
-lmkl_sequential;-lpthread;-lm;-ldl"
$ make
$ make install
```

9. Configuring and Building STARS-H Library (optional),

```
$ cd exageostat
$ cd stars-h
$ mkdir build
$ cmake .. -DCMAKE_C_FLAGS=-fPIC
$ make -j && make install
```

10. Configuring and Building HiCMA library (optional):

```
$ cd ..
# mkdir build && cd build
$ cmake .. -DCMAKE_C_FLAGS=-fPIC -DCBLAS_DIR="${MKLR00T}"
-DLAPACK_DIR="${MKLR00T}"
-DBLAS_LIBRARIES="-L${MKLR00T}/lib;-lmkl_intel_lp64;-lmkl_core;
-lmkl_sequential;-lpthread;-lm;-ldl"
$ make
$ make install
```

Note: Here, we assume ROOT access during installation. if not,

1. `-DCMAKE_INSTALL_PREFIX=$PWD/installdir` should be added to the `cmake` commands, where `installdir` is your local installation directory.
2. `-prefix=$PWD/installdir` should be added to the `configure` commands, where `installdir` is your local installation directory.
3. `pkgconfig` for `NLOPT`, `GSL`, `hwloc`, `StarPU`, `Chameleon`, `STARS-H`, `HiCMA` paths need to be exported to `PKG_CONFIG_PATH` environment variable.

2.3 ExaGeoStat on R

1. Install latest ExaGeoStat R version hosted on GitHub:

```
install.packages("devtools")
library(devtools)
install_git(url="https://github.com/ecrc/exageostatR")
library(exageostat)
```

2.4 Verifying the Installation

Run one or more examples from Chapter 4.

3 CHAPTER 3. LINUX

3.1 Supported Linux Distributions

This chapter provides instructions for installing ExaGeoStat on selected Linux distributions:

1. Ubuntu 16.04 LTS.
2. Fedora Core 25.
3. Red Hat Enterprise Linux Desktop Workstation 7.x
4. OpenSUSE 42.

3.2 Installing the prerequisite Packages

1. Downloading, Configuring and Building CMake 3.2.3 or higher,


```
$ wget https://cmake.org/files/v3.10/cmake-3.10.0-rc3.tar.gz
$ tar -zxvf cmake-3.10.0-rc3.tar.gz
$ cd cmake-3.10.0-rc3
$ ./configure
$ make -j && make -j install
```
2. Download an Install Intel MKL library from <https://software.intel.com/en-us/mkl> for Linux.


```
$ export MKLROOT=/opt/intel/mkl
```
3. Downloading, Configuring and Building NLOPT library 2.4.2 or higher,


```
$ wget http://ab-initio.mit.edu/nlopt/nlopt-2.4.2.tar.gz
$ tar -zxvf nlopt-2.4.2.tar.gz
$ cd nlopt-2.4.2
$ ./configure --enable-shared --without-guile
$ make -j && make -j install
```
4. Downloading, Configuring and Building GSL library 2.4 or higher,


```
$ wget https://ftp.gnu.org/gnu/gsl/gsl-2.4.tar.gz
$ tar -zxvf gsl-2.4.tar.gz
$ cd gsl-2.4
$ ./configure
$ make -j && make -j install
```

5. Downloading, Configuring and Building hwloc 1.11.5 or higher,

```
$ wget https://www.open-mpi.org/software/hwloc/v1.11/
downloads/hwloc-1.11.5.tar.gz
$ tar -zxvf hwloc-1.11.5.tar.gz
$ cd hwloc-1.11.5
$ ./configure
$ make -j && make -j install
```

6. Downloading, Configuring and Building StarPU 1.2.5 or higher,

```
$ wget http://starpup.gforge.inria.fr/files/starpup-1.2.5/
starpup-1.2.5.tar.gz
$ tar -zxvf starpu-1.2.5.tar.gz
$ cd starpu-1.2.5
$ ./configure --disable-cuda --disable-mpi --disable-opencl
$ make -j && make -j install
```

StarPU Configuration: In the case of GPUs systems, both CUDA and OpenCL should be enabled using `-enable` option. If MPI is required, it should be also enabled using `-enable` option.

7. Download Chameleon, HiCMA, and STARS-H,

```
$ git clone https://github.com/ecrc/exageostat.git
$ cd exageostat
$ git submodule update --init --recursive
```

8. Configuring and Building Chameleon Software,

```
$ cd exageostat
$ cd hicma
$ cd chameleon
$ mkdir build && cd build
$ cmake .. -DCHAMELEON_USE_MPI=OFF -DBUILD_SHARED_LIBS=ON
-DCLAS_DIR="${MKLR00T}" -DLAPACK_DIR="${MKLR00T}"
-DCLAS_LIBRARIES="-L${MKLR00T}/lib;-lmkl_intel_lp64;-lmkl_core;
-lmkl_sequential;-lpthread;-lm;-ldl"
$ make
$ make install
```

Chameleon Configuration: In the case of GPUs systems, `-DCHAMELEON_USE_CUDA=ON` If MPI is required, it should be also enabled using `-DCHAMELEON_USE_MPI=ON`.

9. Configuring and Building STARS-H Library (optional),

```
$ cd exageostat
$ cd stars-h
$ mkdir build
$ cmake .. -DCMAKE_C_FLAGS=-fPIC
$ make -j && make install
```

10. Configuring and Building HiCMA library (optional):

```

$ cd ..
$ mkdir build && cd build
$ cmake .. -DCMAKE_C_FLAGS=-fPIC -DCBLAS_DIR="${MKLR00T}"
-DLAPACK_DIR="${MKLR00T}"
-DBLAS_LIBRARIES="-L${MKLR00T}/lib;-lmkl_intel_lp64;-lmkl_core;
-lmkl_sequential;-lpthread;-lm;-ldl"
$ make
$ make install

```

Note: Here, we assume ROOT access during installation. if not,

1. `-DCMAKE_INSTALL_PREFIX=$PWD/installdir` should be added to the `cmake` commands, where `installdir` is your local installation directory.
2. `-prefix=$PWD/installdir` should be added to the `configure` commands, where `installdir` is your local installation directory.
3. `pkgconfig` for `NLOPT`, `GSL`, `hwloc`, `StarPU`, `Chameleon`, `STARS-H`, `HiCMA` paths need to be exported to `PKG_CONFIG_PATH` environment variable.

3.3 ExaGeoStat on R

1. Install latest ExaGeoStat R version hosted on GitHub:

```

install.packages("devtools")
library(devtools)
install_git(url="https://github.com/ecrc/exageostatR")
library(exageostat)

```

3.4 Verifying the Installation

Run one or more examples from Chapter 4.

4 CHAPTER 4. EXAMPLES

4.1 Example 1: Synthetic generation of Geo-spatial data with MLE exact computation,

```

library("exageostat")           #Load ExaGeoStat-R lib.
seed      = 0                    #Initial seed to generate XY locs.
theta1    = 1                    #Initial variance.
theta2    = 0.1                  #Initial range.
theta3    = 0.5                  #Initial smoothness.
dmetric   = 0                    #0 --> Euclidean distance, 1--> great circle distance.
n         = 1600                 #n*n locations grid.
ncores    = 2                    #Number of underlying CPUs.
gpu       = 0                    #Number of underlying GPUs.
ts       = 320                  #Tile_size: changing it can improve the performance.
p_grid   = 1                    #More than 1 in the case of distributed systems.
q_grid   = 1                    #More than 1 in the case of distributed systems.
clb      = vector(mode="double",length = 3) #Optimization lower bounds values.
cub      = vector(mode="double",length = 3) #Optimization upper bounds values.
theta_out  = vector(mode="double",length = 3) #Parameter vector output.
globalveclen = 3*n
vecs_out   = vector(mode="double",length = globalveclen) #Z measurements of n locations.
clb       = as.double(c("0.01", "0.01", "0.01")) #Optimization lower bounds.
cub       = as.double(c("5.00", "5.00", "5.00")) #Optimization upper bounds.
vecs_out[1:globalveclen] = -1.99
theta_out[1:3] = -1.99
exageostat_initR(ncores, gpu, ts)           #Initiate exageostat instance.
vecs_out      = exageostat_egenzR(n, ncores, gpu, ts, p_grid, q_grid,
theta1, theta2, theta3, dmetric, seed, globalveclen) #Generate Z observation vector.
theta_out     = exageostat_emleR(n, ncores, gpu, ts, p_grid, q_grid,
vecs_out[1:n], vecs_out[n+1:(2*n)],
vecs_out[(2*n+1):(3*n)], clb, cub, dmetric, 0.0001, 20) #Exact Estimation (MLE).
exageostat_finalizeR()                       #Finalize exageostat instance
=====

```

4.2 Example 2: Synthetic generation of Geo-spatial data with MLE TLR computation,

```

library("exageostat")           #Load ExaGeoStat-R lib.
seed      = 0                    #Initial seed to generate XY locs.
theta1    = 1                    #Initial variance.
theta2    = 0.03                 #Initial range.
theta3    = 0.5                  #Initial smoothness.
dmetric   = 0                    #0 --> Euclidean distance, 1--> great circle distance.
n         = 900                  #n*n locations grid.
ncores    = 4                    #Number of underlying CPUs.
gpu       = 0                    #Number of underlying GPUs.
ts       = 320                  #Tile_size: changing it can improve the performance.
lts      = 600                  #TLR_Tile_size: changing it can improve the performance.
tlr_acc  = 7                    #approximation accuracy 10^-(acc).
tlr_maxrank = 450               #Max rank.
p_grid   = 1                    #More than 1 in the case of distributed systems.
q_grid   = 1                    #More than 1 in the case of distributed systems.

```



```

clb      = vector(mode="double",length = 3) #Optimization lower bounds values.
cub      = vector(mode="double",length = 3) #Optimization upper bounds values.
theta_out = vector(mode="double",length = 3) #Parameter vector output.
globalveclen = 3*n
vecs_out = vector(mode="double",length = globalveclen)#Z measurements of n locations.
clb      = as.double(c("0.01", "0.01", "0.01")) #Optimization lower bounds.
cub      = as.double(c("5.00", "5.00", "5.00")) #Optimization upper bounds.
vecs_out[1:globalveclen] = -1.99
theta_out[1:3] = -1.99
exageostat_initR(ncores, gpus, ts) #Initiate exageostat instance.
vecs_out = exageostat_egenzR(n, ncores, gpus, ts, p_grid, q_grid,
theta1, theta2, theta3, dmetric, seed, globalveclen) #Generate Z observation vector.
theta_out = exageostat_tlrmlrR(n, ncores, gpus, lts, p_grid, q_grid,
vecs_out[1:n], vecs_out[n+1:(2*n)], vecs_out[(2*n+1):(3*n)],
clb, cub, tlr_acc, tlr_maxrank, dmetric, 0.0001, 20) #TLR Estimation (MLE).
exageostat_finalizeR() #Finalize exageostat instance.

```

4.3 Example 3: Synthetic generation of Geo-spatial data with MLE DST computation,

```

library("exageostat") #Load ExaGeoStat-R lib.
seed      = 0 #Initial seed to generate XY locs.
theta1    = 1 #Initial variance.
theta2    = 0.1 #Initial range.
theta3    = 0.5 #Initial smoothness.
dmetric   = 0 #0 --> Euclidean distance, 1--> great circle distance.
n         = 1600 #n*n locations grid.
ncores    = 2 #Number of underlying CPUs.
gpus      = 0 #Number of underlying GPUs.
ts        = 320 #Tile_size: changing it can improve the performance.
p_grid    = 1 #More than 1 in the case of distributed systems.
q_grid    = 1 #More than 1 in the case of distributed systems.
clb      = vector(mode="double",length = 3) #Optimization lower bounds values.
cub      = vector(mode="double",length = 3) #Optimization upper bounds values.
theta_out = vector(mode="double",length = 3) #Parameter vector output.
globalveclen = 3*n
vecs_out = vector(mode="double",length = globalveclen)#Z measurements of n locations.
clb      = as.double(c("0.01", "0.01", "0.01")) #Optimization lower bounds.
cub      = as.double(c("5.00", "5.00", "5.00")) #Optimization upper bounds.
vecs_out[1:globalveclen] = -1.99
theta_out[1:3] = -1.99
exageostat_initR(ncores, gpus, ts) #Initiate exageostat instance.
vecs_out = exageostat_egenzR(n, ncores, gpus, ts, p_grid, q_grid,
theta1, theta2, theta3, dmetric, seed, globalveclen) #Generate Z observation vector.
theta_out = exageostat_dstmlrR(n, ncores, gpus, ts, p_grid, q_grid,
vecs_out[1:n], vecs_out[n+1:(2*n)],
vecs_out[(2*n+1):(3*n)], clb, cub, dmetric, 0.0001, 20) #DST Estimation (MLE).
exageostat_finalizeR() #Finalize exageostat instance

```

4.4 Example 4: Synthetic generation of measurements based on given Geo-spatial data locations with MLE Exact computation,

```

library("exageostat")           #Load ExaGeoStat-R lib.
seed      = 0                    #Initial seed to generate XY locs.
theta1    = 1                    #Initial variance.
theta2    = 0.1                  #Initial range.
theta3    = 0.5                  #Initial smoothness.
dmetric   = 0                    #0 --> Euclidean distance, 1--> great circle distance.
n         = 1600                 #n*n locations grid.
ncores    = 2                    #Number of underlying CPUs.
gpu       = 0                    #Number of underlying GPUs.
ts       = 320                  #Tile_size: changing it can improve the performance.
p_grid   = 1                    #More than 1 in the case of distributed systems.
q_grid   = 1                    #More than 1 in the case of distributed systems.
clb      = vector(mode="double",length = 3) #Optimization lower bounds values.
cub      = vector(mode="double",length = 3) #Optimization upper bounds values.
theta_out  = vector(mode="double",length = 3) #Parameter vector output.
globalveclen = 3*n
x          = rnorm(n = globalveclen, mean = 39.74, sd = 25.09)
#X dimension vector of n locations.
y          = rnorm(n = globalveclen, mean = 80.45, sd = 100.19)
#Y dimension vector of n locations.
vecs_out   = vector(mode="double",length = globalveclen)#Z measurements of n locations.
clb       = as.double(c("0.01", "0.01", "0.01")) #Optimization lower bounds.
cub       = as.double(c("5.00", "5.00", "5.00")) #Optimization upper bounds.
vecs_out[1:globalveclen] = -1.99
theta_out[1:3] = -1.99
exageostat_initR(ncores, gpu, ts) #Initiate exageostat instance.
vecs_out      = exageostat_egenz_glr(n, ncores, gpu, ts, p_grid, q_grid,
x, y, theta1, theta2, theta3,
dmetric, globalveclen) #Generate Z observation vector based on given locations.
theta_out     = exageostat_emleR(n, ncores, gpu, ts, p_grid, q_grid,
vecs_out[1:n], vecs_out[n+1:(2*n)],
vecs_out[(2*n+1):(3*n)], clb, cub, dmetric, 0.0001, 20) #Exact Estimation (MLE).
exageostat_finalizeR() #Finalize exageostat instance

```